

A young girl with dark hair, wearing a bright pink hooded raincoat, is looking upwards and to the right with a curious expression. She is standing in a dark, cluttered street at night, illuminated by vibrant neon lights in shades of blue, green, and orange. The background is out of focus, showing various street elements and more neon signs.

# Deep Learning for Visual Computing

## Modern Classification Networks

Christopher Pramerdorfer  
Computer Vision Lab, TU Wien

Image from [twitter/mechanicmind\\_ai](https://twitter.com/mechanicmind_ai), created using midjourney

# Topics

## Optimization vs. Machine Learning

- ▶ Regularization

## Modern classification networks

- ▶ Residual networks
- ▶ Efficient architectures

## Pushing classification performance

- ▶ Fine-tuning
- ▶ Handling imbalanced data

# Optimization vs. Machine Learning

## Regularization

Recall from last lecture that our goal is to

- ▶ Reach high performance on unseen (validation/test) data
- ▶ By training on training data

For optimal results

- ▶ Networks should have enough capacity to overfit
- ▶ But be configured not to via regularization

# Optimization vs. Machine Learning

## Regularization

The purpose of **regularization** is to

- ▶ Improve the validation/test performance
- ▶ At the possible expense of training performance

Usually by decreasing the models **variance**

- ▶ Sensitivity to small changes in training set
- ▶ And thus proneness to overfitting

# Optimization vs. Machine Learning

## Regularization – Penalizing Large Weights

A penalty on large weights is almost always used

- ▶ Prevents certain inputs from dominating output
- ▶ Encourages model to use all inputs

Biases are not critical

- ▶ Usually not subject to regularization

# Optimization vs. Machine Learning

## Regularization – L2 Regularization

A common method is **L2 regularization**

- ▶ Implemented by adding regularization term to loss function

$$L_{reg}(\boldsymbol{\theta}) = L(\boldsymbol{\theta}) + \frac{\delta}{2} \|\mathbf{w}\|^2$$

$\mathbf{w} \subset \boldsymbol{\theta}$  is vector of all weights

$\delta \in (0, 1)$  controls amount of regularization

- ▶ Usually  $\delta \in [0.0001, 0.01]$

# Optimization vs. Machine Learning

## Regularization – L2 Regularization

Ignoring bias we get  $\nabla L_{reg}(\mathbf{w}) = \nabla L(\mathbf{w}) + \delta \mathbf{w}$

- ▶  $\|\mathbf{w}\|^2 = \mathbf{w} \cdot \mathbf{w}$  so gradient is  $2\mathbf{w}$  (product rule)

So gradient descent update becomes

$$\begin{aligned}\mathbf{w} &= \mathbf{w} - \alpha(\nabla L(\mathbf{w}) + \delta \mathbf{w}) \\ &= \mathbf{w} - \alpha \delta \mathbf{w} - \alpha \nabla L(\mathbf{w}) \\ &= (1 - \alpha \delta) \mathbf{w} - \alpha \nabla L(\mathbf{w})\end{aligned}$$

# Optimization vs. Machine Learning

## Regularization – L2 Regularization

Weights shrink by constant factor before each update

- ▶ So if  $\nabla L(\mathbf{w}) = \mathbf{0}$ , weights would decay towards  $\mathbf{0}$
- ▶ Note that  $\alpha$  also affects regularization

Decay strength  $\delta$  is another hyperparameter

- ▶ No effect if too small
- ▶ Dominates data loss (e.g. cross-entropy) if too large



# Optimization vs. Machine Learning

## Regularization – Weight Decay

A very similar approach is **weight decay**

- ▶ Do not modify loss function
- ▶ Adapt weight update rule to subtract  $\alpha\delta\mathbf{w}$

L2 regularization affects optimizer

- ▶ Leads to undesirable behavior with e.g. Adam
- ▶ So always use version of Adam with weight decay
- ▶ PyTorch: `torch.optim.AdamW`

# Optimization vs. Machine Learning

## Regularization – Dropout

**Dropout** [4] is a layer whose neurons

- ▶ Output 0 with probability  $p$
- ▶ Forward input unchanged with probability  $1 - p$
- ▶ But only during training


Usually placed before last (dense) layer

- ▶ Has effect of temporarily “dropping” neurons
- ▶ Because 0 does not change output


Not as popular anymore

# Optimization vs. Machine Learning

## Regularization – Dropout



(a) Standard Neural Net



(b) After applying dropout.


Image from [4]

# Optimization vs. Machine Learning

## Early Stopping

Early stopping aims to avoid overfitting by

- ▶ Storing a copy of  $\theta$  with best validation performance  $p$
- ▶ Stopping if  $p$  does not improve anymore for  $e$  epochs
- ▶ Using the stored  $\theta$  afterwards



# Optimization vs. Machine Learning

## Suggestions

### Data

- ▶ Use as much as you can get
- ▶ Always utilize (sensible) training data augmentation

### Regularization

- ▶ Always use weight decay and tune  $\delta$
- ▶ Batch normalization also acts as a regularizer
- ▶ Consider dropout if network still overfits

Always use early stopping and with  $e > 1$




# Modern Classification Networks

## Baseline

Basic design from previous lecture

- ▶ Depth of 10 (number of conv and linear layers)
- ▶ Not deep enough for optimal performance



# Modern Classification Networks

## Training Deep Networks

Increasing depth is simple in theory

- ▶ Just add more conv layers between pooling layers

In practice as a rule of thumb

- ▶ Increasing the depth like this improves performance
- ▶ With batch normalization and proper regularization
- ▶ Up to a depth of around 20



# Modern Classification Networks

## Training Deep Networks

Even deeper networks eventually perform worse

- ▶ Not intuitive (network could learn identity functions)
- ▶ However doing so is challenging






Image from [1]

# Modern Classification Networks

## Residual Networks

Residual networks (ResNets) [1] facilitate this task

- ▶ Learn what to change about input
- ▶ Making it easy to learn identity function
- ▶ By using skip-connections and summation



# Modern Classification Networks

## Residual Networks

Idea is that

- ▶ Individual layers only change input a little
- ▶ But network has many more of them





Image adapted from [1]

# Modern Classification Networks

## Residual Networks

Another benefit is easier gradient flow

- ▶ Recall that sum nodes propagate gradients without changes
- ▶ And that gradients that join are summed

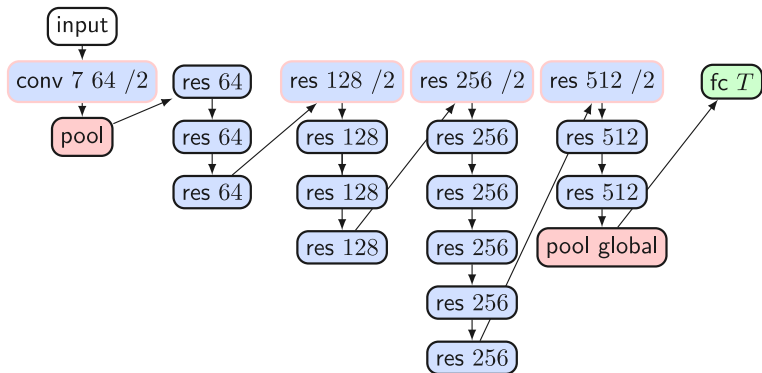


# Modern Classification Networks

## Residual Networks

ResNet-34 ( $R = 224$ )

- First and last layers identical to baseline design



# Modern Classification Networks

## Residual Networks

Must support changes to  $C, H, W$

- ▶ To enable pooling and adjusting feature map counts

$1 \times 1$  (**point-wise**) convolutions are a flexible tool

- ▶ Can adjust  $C$  freely
- ▶ Can adjust  $H, W$  via stride

ResNet variant

- ▶ First  $3 \times 3$  conv with stride 2 that doubles  $C$
- ▶ Extra  $1 \times 1$  conv with same configuration before sum

# Modern Classification Networks

## Residual Networks

Such networks scale to depths of 1000 and more (!)

- ▶ 1202 variant overfits below (training error  $\approx 0$ )

method			error (%)
Maxout [10]			9.38
NIN [25]			8.81
DSN [24]			8.22
	# layers	# params	
FitNet [35]	19	2.5M	8.39
Highway [42, 43]	19	2.3M	7.54 (7.72 $\pm$ 0.16)
Highway [42, 43]	32	1.25M	8.80
ResNet	20	0.27M	8.75
ResNet	32	0.46M	7.51
ResNet	44	0.66M	7.17
ResNet	56	0.85M	6.97
ResNet	110	1.7M	<b>6.43</b> (6.61 $\pm$ 0.16)
ResNet	1202	19.4M	7.93


Image from [1]

# Modern Classification Networks

## Residual Networks

Very deep ResNets use more efficient blocks

- Reduce  $C$  before expensive  $3 \times 3$  conv (bottleneck layer)





# Modern Classification Networks

## Residual Networks

ResNets still perform great today

- ▶ Accuracy typically within  $\approx 5\%$  of state of the art
- ▶ Good default architecture for deep networks
- ▶ PyTorch: `net = torchvision.models.resnet34()`

Not optimized for efficiency though

# Modern Classification Networks

## Efficient Architectures

Efficiency usually matters

- ▶ More efficient networks cost less to operate
- ▶ Hardware limitations (e.g. mobile phones)

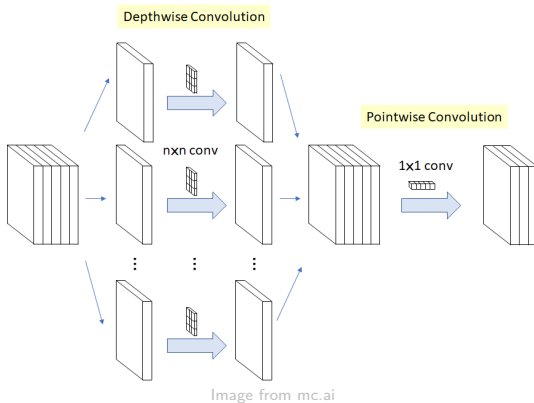
We focus on efficiency during **inference** (not training)

- ▶ In terms of accuracy vs. FLOPS (or predictions/second)

# Modern Classification Networks

## Efficient Architectures

Point-wise and **depth-wise** convolutions are key ingredients



# Modern Classification Networks

## Efficient Architectures

### Building blocks of MobileNet v2 [3]





Image from [2]

# Modern Classification Networks

## Efficient Architectures



# Modern Classification Networks

## Network Scaling

We previously talked about varying the depth of our networks

This is a form of **network scaling**

- ▶ Adapt capacity of network
- ▶ To optimize performance given task at hand
- ▶ While considering the computational budget

We will next cover other forms

# Modern Classification Networks

## Network Scaling





Image from [3]

# Modern Classification Networks

## Network Scaling

Scaling width, depth, and resolution are all effective

- ▶ Impact on FLOPS varies (depends on architecture)
- ▶ Improvements compound (with diminishing returns)



Scaling by width, depth, resolution, in this order. Image from [3].



# Modern Classification Networks

## Network Scaling

Given all this, how can we design efficient networks?

For practitioners like you

- ▶ Stick to established architectures that work well
- ▶ ResNet, EfficientNet, ConvNeXt are examples

For researchers in the field


- ▶ Manually via trial and error (based on prior work)
- ▶ Automatically via **neural architecture search (NAS)**

# Modern Classification Networks

## ConvNeXt

State of the art architecture


- Design inspired by Transformers (later)



# Modern Classification Networks

## ConvNeXt

**ResNet Block**



**ConvNeXt Block**





Image from [5]

# Modern Classification Networks

## ConvNeXt





# Pushing Classification Performance

## Fine-Tuning

Data augmentation & regularization cannot replace data

- ▶ Deep learning scales extremely well with data
- ▶ So obtaining more data should be a priority

For this reason, a powerful technique is **fine-tuning**

- ▶ **Pre-train** network on related large dataset (e.g. ImageNet)
- ▶ Then fine-tune same network on smaller target dataset

Idea is to exploit information present in both datasets

# Pushing Classification Performance

## Fine-Tuning

For instance we could

- ▶ Pre-train a network on ImageNet (1000 classes)
- ▶ Then replace its linear classifier with one for 2 classes
- ▶ And fine-tune the network on our cats vs. dogs dataset

Fine-tuning almost always helps

- ▶ Even if target dataset is large itself
- ▶ Assuming both datasets are related (the more, the better)

# Pushing Classification Performance

## Fine-Tuning

If the target dataset is very small

- ▶ Consider fine-tuning just the new classifier
- ▶ Can be done by freezing parameters of other layers

If not, a good strategy is to

- ▶ First fine-tune only new classifier
- ▶ Then unfreeze other layers and train some more
- ▶ Using a larger learning rate for classifier than for rest



# Pushing Classification Performance

## Handling Imbalanced Data

So far we have assumed balanced data

- ▶ Same number of samples per class

If this is not the case

- ▶ Classifier will pay more attention to majority classes
- ▶ Leading to bad accuracy for minority classes

This is usually not desired and must be addressed

- ▶ We will cover two popular options

# Pushing Classification Performance

## Handling Imbalanced Data

Via **class weights** on the loss

- ▶ Multiply per-sample losses by weights
- ▶ With weights based on inverse class frequencies
- ▶ Easy to implement and no computational overhead

Via **oversampling** of the training data

- ▶ Draw samples with replacement to balance class frequencies
- ▶ Works well in combination with data augmentation
- ▶ Considerable computational overhead

We will now move on from classification to other tasks

- ▶ Most architectures use a classification network as **backbone**
- ▶ Train a classification network on large dataset (e.g. ImageNet)
- ▶ Adapt network for e.g. object detection

This is a form of **transfer learning**

- ▶ Adapt a model trained on one task to perform well on another

# Bibliography

- [1] He et al. Deep Residual Learning for Image Recognition. 2015
- [2] Sandler et al. MobileNetV2. 2018
- [3] Tan & Le. EfficientNet. 2019
- [4] Srivastava et al. Dropout. 2014
- [5] Liu et al. A ConvNet for the 2020s. 2022