

A photograph of Snoop Dogg standing in a vast field of purple lavender flowers. He is wearing a purple jacket over a white t-shirt, gold-rimmed glasses, and a gold chain with a cross pendant. He has a slight smile and is looking upwards and to the left. The background shows a soft-focus landscape with green trees and distant mountains under a pale sky.

Deep Learning for Visual Computing

Object Detection

Christopher Pramerdorfer
Computer Vision Lab, TU Wien

Image from reddit, created using midjourney

Topics

Beyond image classification

Object detection

- ▶ Introduction & performance metrics

Deep object detection

- ▶ Sliding window approach
- ▶ Two-stage detectors (R-CNN)
- ▶ One-stage detectors (YOLO)
- ▶ Feature pyramid networks

Beyond Image Classification

We have focused on image classification

- ▶ Fundamental computer vision task
- ▶ We now know how to achieve good performance

Rest of lecture will focus on more challenging tasks

Beyond Image Classification

Recall our ingredients for image classification

- ▶ A suitable dataset
- ▶ A network with a suitable final layer (linear layer)
- ▶ A suitable loss function $L(\theta)$ (cross-entropy)
- ▶ An algorithm for computing $\nabla L(\theta)$ (backpropagation)
- ▶ An algorithm for updating θ on this basis (Adam)
- ▶ An intuitive performance metric (accuracy)

Beyond Image Classification

These ingredients are universal

Steps 4 and 5 are not task-specific and can be reused

So when tackling a new problem, ask yourself

- ▶ Do I have enough (labeled) data?
- ▶ What is a suitable network output?
- ▶ What is a suitable loss function?
- ▶ What is a good metric for evaluation?

Beyond Image Classification

Answer to first question is almost always “no”

- ▶ Networks are usually pre-trained on ImageNet
- ▶ First learn what different objects look like
- ▶ Then adapt for other tasks (transfer learning)

PyTorch provides such models directly

- ▶ `net = models.resnet34(pretrained=True)`

Object Detection

Challenges & Opportunities

More challenging than image classification

- ▶ Same basic challenges (see lecture 1)
- ▶ More complex task
- ▶ Harder to implement efficiently

Many useful applications such as

- ▶ Face detection (e.g. surveillance)
- ▶ Autonomous driving (e.g. road sign detection)

Object Detection

Datasets

COCO is the most popular dataset with detection labels



Image from cocodataset.org

Object Detection

Progress




Image from paperswithcode.com

What is a suitable object detector output?

Most detectors predict the following

- ▶ N bounding boxes \mathcal{B}_n
- ▶ N corresponding class labels c_n
- ▶ N corresponding confidence scores p_n

Object Detection

Loss Function

Two tasks at once

- ▶ Predict accurate bounding boxes (regression)
- ▶ Assign them correct class labels (classification)

We thus usually have $L(\theta) = L_{\text{loc}}(\theta) + L_{\text{clf}}(\theta)$

- ▶ First term measures bounding box accuracy (e.g. L1 loss)
- ▶ Second term is classification loss (e.g. cross-entropy loss)

Object Detection

Performance Metrics

Intersection over union (IoU) base measure


$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$


Image from Wikipedia

Object Detection

Performance Metrics

Ignoring class labels and given a

- ▶ Single ground-truth bounding box \mathcal{B}'
- ▶ Single predicted bounding box \mathcal{B}
- ▶ IoU threshold t_{iou}

\mathcal{B} is a

- ▶ True positive (TP) if $\text{iou}(\mathcal{B}', \mathcal{B}) \geq t_{iou}$
- ▶ False positive (FP) otherwise

In second case there is also one false negative (FN)

- ▶ No \mathcal{B} with sufficient IoU with \mathcal{B}'

Object Detection

Performance Metrics

Ignoring class labels and given multiple bounding boxes

- ▶ **Precision** = $TPs / (TPs + FPs)$
- ▶ **Recall** = $TPs / (TPs + FNs)$

In case of multiple \mathcal{B} for a given \mathcal{B}'

- ▶ Count \mathcal{B} with highest p as TP, others als FPs

Both metrics are in $[0, 1]$

- ▶ Ideally precision = recall = 1

Object Detection

Performance Metrics

Use confidence threshold t_p to balance recall vs. precision

Increasing t_p results in fewer detections


- ▶ Usually increases precision
- ▶ Usually decreases recall

Object Detection

Performance Metrics

Precision vs. recall curves highlight this behavior

- ▶ Plotting precision and recall over $t_p \in [0, 1]$
- ▶ Area under curve is called **average precision** (AP)



Object Detection

Performance Metrics

AP is most popular base metric

- ▶ Concise measure of detector performance
- ▶ For single t_{iou} and single class

Mean average precision (mAP) extends this to C classes

$$\text{mAP}_{t_{iou}} = \frac{1}{C} \sum_c^C \text{AP}_c$$

And further to multiple t_{iou} , e.g.

$$\text{mAP} = \frac{1}{10}(\text{mAP}_{0.50} + \dots + \text{mAP}_{0.95})$$

Confusingly mAP is sometimes named just AP in papers

- ▶ Like in COCO performance chart above

Deep Object Detection

Many interesting approaches and methods exist

- ▶ See survey papers such as [1] for an overview

Virtually all use a classification network as **backbone**

- ▶ Transfer learning as mentioned earlier

We will cover three approaches

- ▶ Naive sliding window approach
- ▶ Two-stage detectors (R-CNN)
- ▶ Single-stage detectors (YOLO)

Deep Object Detection

Sliding Window Approach

Training

- ▶ Train a classifier for $C + 1$ classes
- ▶ Additional “background” class

Detection

- ▶ Slide fixed-size window over image
- ▶ Predict class-scores for every window
- ▶ Perform non-maximum suppression

Deep Object Detection

Sliding Window Approach




Image adapted from apple.com

Deep Object Detection

Sliding Window Approach

Inefficient

- ▶ Many windows to classify

Single fixed-size window (no scale invariance)

- ▶ Must process image at multiple scales (more inefficient)
- ▶ Inaccurate bounding boxes (fixed aspect ratio)

Cannot handle multiple objects in same window (softmax)

Deep Object Detection

Two-Stage Detectors

Improve efficiency by

- ▶ First generating many **region proposals**
- ▶ Classifying these proposals




Image adapted from [1]

Deep Object Detection

Two-Stage Detectors

Region proposals have different size and aspect ratio

- ▶ Solves window-related problems
- ▶ Warp to common size to support minibatch predictions





Image adapted from [1]

Deep Object Detection

Two-Stage Detectors

Approach of R-CNN

- ▶ First successful deep-learning-based object detector
- ▶ Clunky design and training, see [2] if interested



Deep Object Detection

Two-Stage Detectors

Approach still inefficient

- ▶ Got to classify many region proposals

More efficient extension

- ▶ Process complete (high-resolution) image once
- ▶ Project region proposals onto suitable conv layer output
- ▶ Classify each resulting feature map region

Deep Object Detection

Two-Stage Detectors

Approach was introduced by [Fast R-CNN](#) [3]

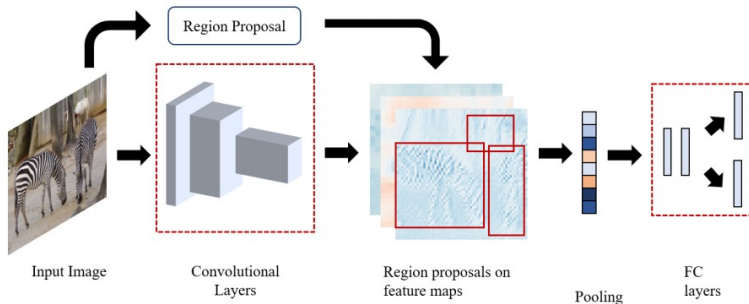


Image adapted from [1]

Deep Object Detection

Two-Stage Detectors

Region proposal step remains a bottleneck

- ▶ Generic algorithms (e.g. selective search [4])
- ▶ Not integrated into network
- ▶ Usually quite slow



Image adapted from [4]

Deep Object Detection

Two-Stage Detectors

There are improvements that address this

- ▶ Faster R-CNN is an example

R-CNN variants are **two-stage detectors**

- ▶ First generate r region proposals
- ▶ Then analyze these proposals

Modern detectors are **one-stage detectors**


- ▶ Combine both stages, inference time independent of r

Deep Object Detection

One-Stage Detectors

Most one-stage detectors utilize **anchor boxes**

- ▶ Fixed set of reference bounding boxes
- ▶ At different scales and aspect ratios



Deep Object Detection

One-Stage Detectors

YOLO is a popular anchor-based detector family

- ▶ High detection performance
- ▶ Runs in real-time

Many versions exist

- ▶ We will cover v2, and v3
- ▶ Overall approach (detector head) identical

Deep Object Detection

YOLO

Simplified overall approach

- ▶ Divide input image into coarse $S \times S$ grid
- ▶ For each grid cell and anchor box, predict class labels (if any)
- ▶ Perform non-maximum suppression




Image from pjreddie.com

Deep Object Detection

YOLO v2

More precisely, given an anchor box we seek to predict

- ▶ The IoU overlap p with the most relevant object
- ▶ The offset and scale of the box (x, y, w, h)
- ▶ C class scores

To constrain the number of detections

- ▶ Divide the image into grid (e.g. 7×7)
- ▶ Align each anchor box with the center of each grid cell
- ▶ Predict the above for every box and every cell

Deep Object Detection

YOLO v2

Assuming k anchor boxes and a 7×7 grid

- ▶ Output is $k \cdot (5 + C) \times 7 \times 7$
- ▶ Can be implemented using conv and pooling layers

Backbone pools down to desired grid size

- ▶ For instance $224 \mapsto 112 \mapsto 56 \mapsto 28 \mapsto 14 \mapsto 7$

Detector head consists of conv layers only

- ▶ Last one has $k \cdot (5 + C)$ channels
- ▶ YOLO v2 uses 4 layers in total (implementation detail)

Deep Object Detection

YOLO v2

Network has no linear layers

- ▶ Fully convolutional network (FCN) architecture
- ▶ Supports varying input (image) resolution

Deep Object Detection

YOLO v2 – Training

Pretrain backbone, then fine-tune for detection

For every training image, grid cell, and bounding box label

- ▶ Find the anchor box with the highest IoU
- ▶ Compute and assign $(p, x, y, h, w, o_1, \dots, o_C)$
- ▶ Assign $(0, \dots)$ to all unassigned anchor boxes

Deep Object Detection

YOLO v2 – Training

Weighted sum of three L2 losses per anchor box

- ▶ Confidence loss (loss on p)
- ▶ Classification loss (loss on class scores) *
- ▶ Localization loss (loss on x, y, w, h) *

(*) Only if IoU with ground-truth box is sufficient

Deep Object Detection

YOLO v2 – Inference

Network always predicts $k \cdot S \cdot S$ detections

- ▶ But most have $p \approx 0$
- ▶ Perform per-class non-maximum suppression



Image adapted from [Youtube](#)

Deep Object Detection

YOLO v2 – Limitations

Issues with

- ▶ Small objects
- ▶ Predicting tight bounding boxes

Because backbone feature tensor has low resolution of $S \times S$


- ▶ Late in network so semantically strong features
- ▶ But limited amount of spatial information due to pooling

Deep Object Detection

Feature Pyramid Networks

Could address this by duplicating detector heads

- ▶ Attach to different conv layers



(c) Pyramidal feature hierarchy


Image from [6]

Deep Object Detection

Feature Pyramid Networks

Earlier layers have higher resolution but are semantically weaker

- We want both



(c) Pyramidal feature hierarchy


Image from [6]

Deep Object Detection


Feature Pyramid Networks

Feature pyramid networks (FPNs) achieve this by combining

- ▶ High-resolution but semantically weaker features
- ▶ Lower-resolution but semantically stronger features



(c) Pyramidal feature hierarchy



(d) Feature Pyramid Network


Image from [6]

Deep Object Detection

Feature Pyramid Networks

Accomplished by

- ▶ Upsampling lower-resolution features
- ▶ Combining features using sum operation




Deep Object Detection

Feature Pyramid Networks

Requires matching number of channels

- ▶ Fixed number of channels (e.g. 256)
- ▶ Implemented using pointwise convolutions



Deep Object Detection

Feature Pyramid Networks

FPNs aggregate and refine features

- ▶ Independent of base network architecture and task

In object detection context

- ▶ FPN (or alternatives) forms the **neck** of the detector
- ▶ Independent of backbone and head network architectures

Deep Object Detection

YOLO v3

YOLO v3 uses a FPN neck with 3 scales

- ▶ One YOLO head per scale
- ▶ In addition to other improvements, see [7] if interested




Image adapted from [7]

Several improved variants of YOLO exist, such as [8]

- ▶ Better backbone network
- ▶ Improved neck architecture
- ▶ Better training data augmentation
- ▶ Overall approach (detector head) largely identical

YOLO is a good candidate for object detection

- ▶ Modern variants perform well (COCO mAP around 55%)
- ▶ Have great speed vs. performance trade-off

Deep Object Detection

YOLO – Remarks

Speed vs. performance on COCO [8]




Image from [8]

Deep Object Detection

YOLO – Showcase



Image from [youtube](#)

Bibliography

- [1] Zaidi et al. [Deep Object Detection Survey](#). 2021
- [2] Girshick et al. [R-CNN](#). 2013
- [3] Girshick. [Fast R-CNN](#). 2015
- [4] Uijlings et a. [Selective Search](#). 2013
- [5] Redmon & Farhadi. [YOLO v2](#). 2016
- [6] Lin et al. [Feature Pyramid Networks](#). 2017
- [7] Redmon & Farhadi. [YOLO v3](#). 2017
- [8] Wang et al. [Scaled-YOLO v4](#). 2021