



Deep Learning for Visual Computing

Dense Prediction

Christopher Pramerdorfer
Computer Vision Lab, TU Wien

Topics

Dense prediction

- ▶ Semantic image segmentation
- ▶ Keypoint detection
- ▶ Image restoration

Building blocks

- ▶ Upsampling layers
- ▶ U-Nets & Conditional GANs

This Week in AI

Drag Your GAN



Image from [youtube](#)

Dense Prediction

We often want to predict something for every input image pixel

- ▶ This is called **dense prediction**
- ▶ How can we achieve that?


Obvious solution is to avoid pooling

- ▶ We already established this is too slow
- ▶ Also receptive field increases slowly so missing context

Dense Prediction

Solution is to use an **encoder-decoder architecture**


- **Encoder** learns high-level features
- **Decoder** maps these features back to original size



Dense Prediction

Layer types

- ▶ Encoder: conv and pooling layers (e.g. ResNet)
- ▶ Decoder: conv and **upsampling** layers



Dense Prediction

Upsampling – Interpolatuion

One way is to utilize standard 2D upsampling techniques

- ▶ Nearest neighbor or bilinear interpolation
- ▶ Independently per channel

Fast and simple

Usually paired with a conv layer

- ▶ To learn useful features / transformations

Dense Prediction

Upsampling – Interpolation

Example (“inverse max-pooling”)

- ▶ Upsampling by factor of 2
- ▶ Nearest neighbor interpolation

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \Rightarrow \begin{bmatrix} 1 & 1 & 2 & 2 \\ 1 & 1 & 2 & 2 \\ 3 & 3 & 4 & 4 \\ 3 & 3 & 4 & 4 \end{bmatrix}$$

Dense Prediction

Upsampling – Transposed Convolutions

Transposed convolutions are convolutions with stride $1/s$

- ▶ Also called **fractionally strided convolutions**
- ▶ See link below for animations of variants




Image from [github](#)

Dense Prediction

Upsampling – Transposed Convolutions

Learnable upsampling

- ▶ More powerful than pure interpolation
- ▶ But often leads to artifacts due to overlapping kernel




Image from [github](#)

Dense Prediction

Upsampling – Subpixel Convolutions

Subpixel convolutions are most recent upsampling method

- ▶ Usually outperforms other methods covered
- ▶ Also called pixel shuffle

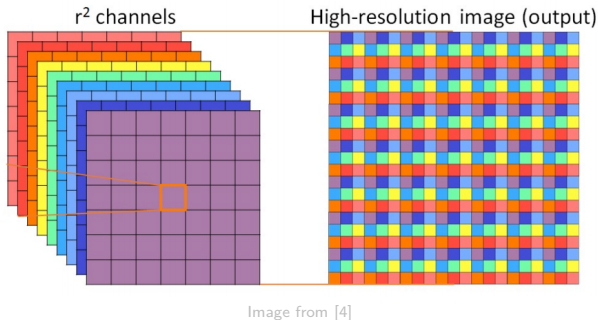
Operation

- ▶ Given input of size $(C \cdot r^2) \times H \times W$
- ▶ Rearrange (shuffle) to $C \times rH \times rW$
- ▶ Then do a stride 1 convolution

Dense Prediction

Upsampling – Subpixel Convolutions


Example with $r = 3$ and $9 \times 7 \times 7$ input



Dense Prediction

We can train such networks like we did before


- ▶ Gradient descent and backpropagation
- ▶ Just define a suitable loss depending on task



Dense Prediction

However this would probably not work very well because


- ▶ The network has no idea how objects look like
- ▶ Spatial information is lost due to excessive downsampling



Dense Prediction

To address this we

- ▶ Pre-train the encoder (e.g. ImageNet classification)
- ▶ Use skip-connections across the two parts




Dense Prediction

U-Nets

A popular method for defining these connections is [U-Net](#) [2]

- ▶ Skip from before every downsampling layer
- ▶ To after every corresponding upsampling layer
- ▶ [Concatenate](#) (not sum) signals



Dense Prediction

U-Nets

Skip-connections forward high-resolution data

- ▶ Makes spatial information available to decoder
- ▶ Signals are concatenated due to different nature




Image from medium.com

Dense Prediction

U-Nets

Modern U-Nets often use

- ▶ A ResNet for downsampling
- ▶ Subpixel convolutions for upsampling




Image from medium.com

Dense Prediction

Applications

We now have all we need to solve various tasks

- ▶ Simply choose C of output accordingly
- ▶ And use a suitable loss function
- ▶ Always pre-train the encoder

We will cover three popular applications

Semantic Image Segmentation

Task definition

- ▶ Given T classes
- ▶ Assign a class label to each image pixel



Image from cocodataset.org

Semantic Image Segmentation

Implementation

- ▶ Configure last layer to output T channels
- ▶ Use cross-entropy loss $L_p(\boldsymbol{\theta})$ for every pixel p
- ▶ Minimize $L(\boldsymbol{\theta}) = \sum_p L_p(\boldsymbol{\theta})$

Keypoint Detection

Keypoint detection task definition

- ▶ Given K keypoints
- ▶ Locate each keypoint in image (multiple times)

A popular application is (sparse) human pose estimation

- ▶ Keypoints for hands, feet, shoulders etc.

Keypoint Detection

Implementation

- ▶ Configure last layer to output K channels
- ▶ Each channel encodes keypoint confidences (**confidence map**)
- ▶ Use e.g. L2 loss $L_c(\theta)$ (**MSE loss**) for every channel c
- ▶ Minimize $L(\theta) = \sum_c L_c(\theta)$

Supporting multiple instances requires keypoint matching

- ▶ Can be done in post-processing step

Keypoint Detection

Note that dense prediction is not mandatory

- ▶ Could simply regress keypoint coordinates

But dense prediction is more flexible & powerful

- ▶ Confidences facilitate downstream tasks
- ▶ Can predict additional dense data (e.g. part affinity fields)




Image from [3]

Keypoint Detection

Human pose estimation with OpenPose [3]



Source: <https://www.youtube.com/watch?v=2DlOUK11YaY>

[link](https://www.youtube.com/watch?v=2DlOUK11YaY)

Image Restoration

Image restoration aims to improve image quality

Common tasks

- ▶ Remove overlays and compression artifacts
- ▶ Sharpen or upscale




Image from medium.com

Image Restoration

Implementation

- ▶ Configure last layer to output C_0 channels (usually $C_0 = 3$)
- ▶ Train on pairs of low-quality and high-quality images
- ▶ Use L2 loss $L_p(\theta)$ for every pixel p
- ▶ Minimize $L(\theta) = \sum_p L_p(\theta)$

Image Restoration

Generating large datasets is comparatively easy

- ▶ Take high-quality images
- ▶ Degrade them automatically depending on task (e.g. blur)




Image from medium.com

Image Restoration

The resulting network would probably be

- ▶ Good at removing text and compression artifacts
- ▶ Not so good at sharpening the image



Image from medium.com

Image Restoration

Because L2 loss does not capture sharpness well

- ▶ Blurring has little effect on most pixel values
- ▶ Per-pixel losses cannot capture overall sharpness



Image from [4]

Image Restoration

Unclear how to define a good loss for sharpness

- ▶ Maybe we can “learn” it?

To do so we use an **adversarial loss**

- ▶ Loss function is a binary classifier \mathcal{D} (**discriminator** or **critic**)
- ▶ That answers “is the given image real or fake?”

In our context of image restoration

- ▶ “real” means an original high-quality image
- ▶ “fake” means one processed by the network

Image Restoration

GANs

This is a **Conditional Generative Adversarial Network (GAN)**

- ▶ **Generator** \mathcal{G} (our encoder-decoder) learns to fool the critic
- ▶ Critic learns not to be fooled

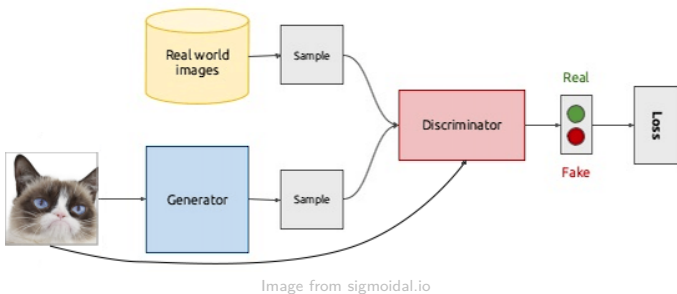


Image Restoration

GANs

If \mathcal{G} can reliably fool \mathcal{D}

- ▶ Processed images look just like high-quality ones
- ▶ And thus must be high-quality themselves

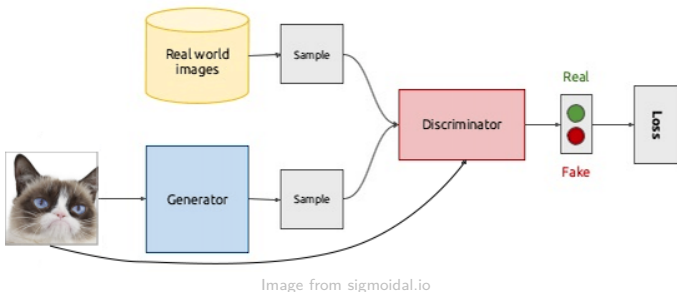


Image Restoration

GANs

\mathcal{G} and \mathcal{D} form a single network during training

- ▶ Gradients flow from \mathcal{D} to \mathcal{G}
- ▶ So \mathcal{G} learns via (cross-entropy) loss on \mathcal{D}

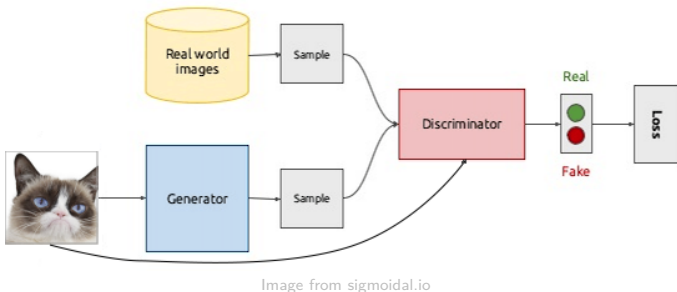


Image Restoration

GANs

Training is an iterative process

- ▶ Train \mathcal{D} (but not \mathcal{G}) on real and fake images
- ▶ Train \mathcal{G} (but not \mathcal{D}) on fake images
- ▶ Repeat

Loss usually does not converge

- ▶ Look at generated images during training
- ▶ Accuracy of \mathcal{D} should eventually be ≈ 0.5

Typical training loop (many variants exist)

- ▶ k is a hyperparameter (often $k = 1$ works fine)

while not done:

for k iterations:

sample half of mini-batch from dataset (class real)

sample half of mini-batch from G (class fake)

train D on mini-batch (but not G)

sample mini-batch from G (class real (!))

train G through D on mini-batch (but not D itself)

GANs are a flexible framework with many applications

- ▶ More in next lecture

Make sure to pre-train both \mathcal{G} and \mathcal{D} (if possible)

- ▶ Training GANs from scratch is a pain

Bibliography

- [1] Noh et al. [Deconvolution Networks](#). 2015
- [2] Ronneberger et al. [U-Net](#). 2015
- [3] Cao et al. [OpenPose](#). 2018
- [4] Shi et al. [Subpixel Convolutions](#). 2016